

Natural Language Processing Laboratory (CSDC-0240)

FINAL PROJECT SUBMISSION

**TITLE: ECHOES OF EMOTIONS (SENTIMENT
ANALYSIS)**

*Name: Harleen Kaur, Anu
Roll no.: 23117012, 23117002
Group: G4A
Branch: Data Science
Submitted to: Ms. Sukhwinder Kaur*



Department of Computer Science and Engineering

DR. B.R. AMBEDKAR NATIONAL INSTITUTE OF
TECHNOLOGY JALANDHAR, PUNJAB-144011

Abstract

This project presents “**Echoes of Emotion**,” a sentiment analysis system designed to evaluate customer feedback using Natural Language Processing (NLP) techniques. The system employs the **VADER (Valence Aware Dictionary and Sentiment Reasoner)** model for rule-based sentiment classification, integrated with a user-friendly **Streamlit** web interface. Key features include **sentence-level sentiment scoring**, **text preprocessing** (tokenization and part-of-speech tagging), and **automated summarization of likes and dislikes** using phrase extraction. The platform processes individual user input efficiently and displays insights in an interactive and accessible manner. By leveraging libraries such as **NLTK** and **Streamlit**, the system ensures reliable performance and usability. This tool provides a practical solution for understanding customer sentiment, with future enhancements planned for batch processing and advanced visualization.

Table of Contents

Abstract	1
1 Introduction	3
2 Literature Survey	3
3 System Architecture	4
4 Tools and Technologies Used	5
5 Implementation	5
5.1 Overview	5
5.2 Core Dependencies and Configuration	5
5.3 Code Implementation	5
5.4 Web Interface	8
6 Results	9
7 Challenges and Solutions	10
7.1 Challenge: Handling Noisy Data	10
7.2 Challenge: Summary Accuracy	10
7.3 Challenge: Interface Responsiveness	10
7.4 Challenge: Error Handling	10
8 Conclusion	10

9	References.....	11
10	Appendix.....	11
10.1	System Requirements.....	11
10.2	Dependencies.....	12

1 Introduction

In the era of digital platforms and customer-centric services, feedback plays a crucial role in shaping business strategies and improving product quality. Manual analysis of user feedback is time-consuming and inefficient, especially when the volume of responses is high. **Sentiment analysis**, a subfield of **Natural Language Processing (NLP)**, offers a practical solution by automatically classifying textual feedback as **positive**, **negative**, or **neutral**, enabling businesses to extract valuable insights.

The "**Echoes of Emotion**" project presents an NLP-based sentiment analysis system that evaluates customer feedback in real time. Unlike systems trained on preloaded datasets, this tool operates on **manually entered input**, using **VADER (Valence Aware Dictionary and Sentiment Reasoner)** to compute sentiment scores and identify prominent **likes and dislikes** using linguistic phrase extraction. The primary objectives of the project are:

- **Analyse Sentiment:** Classify feedback using VADER to determine sentiment polarity and intensity.
- **Summarize Feedback:** Extract frequently mentioned liked and disliked phrases using POS-based chunking and frequency analysis.
- **User-Friendly Interface:** Present the results through an intuitive **Streamlit** web application.
- **Modular Design:** Ensure flexibility for future enhancements such as batch review processing and advanced visualization.

The scope of the project includes **real-time feedback analysis**, **VADER-based sentiment classification**, **noun phrase summarization**, and deployment via a **Streamlit interface**. This report documents the system's architecture, implementation, challenges encountered, and opportunities for future improvements.

2 Literature Survey

Sentiment analysis has seen significant advancements with the evolution of **Natural Language Processing (NLP)**. Rule-based models like **VADER (Valence Aware Dictionary and Sentiment Reasoner)**, introduced by Hutto and Gilbert (2014), are particularly effective for short, informal texts such as customer reviews and social media posts. VADER uses a lexicon-based approach that accounts for **sentiment intensity**, and is capable of interpreting **emoticons, punctuation, and informal language**, making it ideal for analysing real-world feedback data.

Text preprocessing techniques such as **tokenization** and **part-of-speech (POS) tagging** play a vital role in improving the performance of sentiment analysis models. According to Jurafsky and Martin (2023), these methods help in **cleaning text**, **extracting linguistic features**, and enhancing model accuracy by reducing noise.

For deployment, **Streamlit** has emerged as a popular Python framework for building interactive web applications. Its ease of use and lightweight architecture make it well-suited for **deploying NLP applications** to end users, including those without technical expertise.

Recent research also emphasizes the value of **summarizing sentiment output** by extracting frequently mentioned **liked and disliked phrases** from text. Such summarization not only helps in interpreting overall sentiment but also provides **actionable insights** into user preferences — a feature implemented in this project.

3 System Architecture

The “**Echoes of Emotion**” system follows a modular three-tier architecture:

1. **User Interface Layer:** Built with **Streamlit**, this layer provides an interactive web interface where users can input feedback and receive sentiment analysis results. It displays overall sentiment, detailed sentiment scores, and a summary of commonly liked and disliked phrases.
2. **Core Processing Layer:** Implements the main logic using the **CustomerFeedbackAnalyzer** class. This layer handles **text preprocessing** (tokenization and part-of-speech tagging) applies **VADER sentiment analysis** and extracts meaningful noun phrases for summarization. It also handles basic error checking for empty input.
3. **Data Layer:** Currently supports **real-time user input** from the interface. While the architecture is extensible for dataset-based analysis.

Workflow

1. The user enters free-text feedback through the **Streamlit interface**.
2. The **Core Processing Layer** performs tokenization and **POS tagging**.
3. The **VADER model** analyzes each sentence and generates sentiment scores (positive, neutral, negative, and compound).
4. The system identifies and extracts **noun phrases** from positive and negative sentences to summarize **likes and dislikes**.
5. The results, including sentiment score and top phrases, are displayed via the **Streamlit app**, using clearly formatted labels and optional emojis.

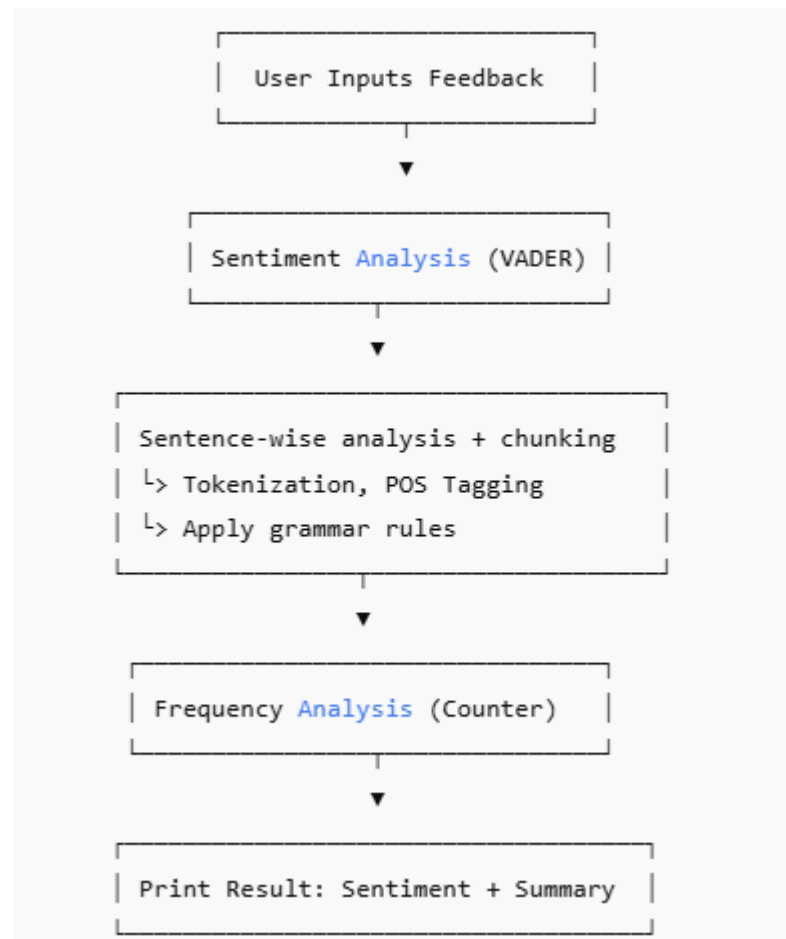


Fig 1: Workflow

Data Flow

The data flow begins with the raw text input, which undergoes preprocessing to remove noise and extract relevant features. Following this, the VADER sentiment analysis model computes sentiment scores, including negative, neutral, positive, and compound metrics.

A custom summarization module then identifies key themes—such as likes and dislikes—based on the sentiment scores and linguistic patterns within the text. Finally, the results are formatted for display, with built-in error handling to manage empty or invalid inputs, ensuring robustness throughout the process.

4 Tools and Technologies Used

The system leverages a robust set of tools and technologies to ensure efficient development and functionality:

Programming Language:

Python is used for its extensive natural language processing (NLP) and data science libraries.

Libraries and Frameworks:

- **NLTK:** Utilized for text preprocessing tasks such as tokenization, part-of-speech (POS) tagging, named entity recognition (NER), and VADER sentiment analysis.
- **Streamlit:** Employed to build the interactive web interface, providing a user-friendly experience.
- **python-dotenv:** Used for managing environment variables securely.

Platforms:

- **Jupyter Notebook:** Used as the primary development environment for iterative coding and testing.
- **GitHub:** Used for version control and collaborative development.

5 Implementation

5.1 Overview

The *Echoes of Emotion* system integrates VADER-based sentiment analysis with a Streamlit interface. It preprocesses the input text, computes detailed sentiment scores, and generates concise summaries. The results are delivered through an interactive web application, enabling users to explore sentiment insights easily.

5.2 Core Dependencies and Configuration

The system requires the following dependencies:

- **NLTK (version 3.8.0 or higher):** Used for text preprocessing and VADER sentiment analysis.
- **Streamlit (version 1.20.0 or higher):** Provides the web interface for user interaction.
- **python-dotenv (version 1.0.0 or higher):** Manages environment variables securely.

Additionally, the VADER lexicon is downloaded through NLTK.

5.3 Code Implementation

```
import nltk
from nltk.sentiment import SentimentIntensityAnalyzer

nltk.download('vader_lexicon')

class CustomerFeedbackAnalyzer:
    def __init__(self):
        """Initialize sentiment analyzer"""
        self.sentiment_analyzer = SentimentIntensityAnalyzer()

    def analyze_sentiment(self, text):
        """Analyze sentiment using NLTK's VADER"""
        scores = self.sentiment_analyzer.polarity_scores(text)
        compound = scores['compound']

        if compound >= 0.05:
            sentiment = 'positive'
        elif compound <= -0.05:
            sentiment = 'negative'
        else:
            sentiment = 'neutral'

        return {
            'sentiment': sentiment,
            'score': compound,
            'details': scores
        }

    def generate_summary(self, text):
        """Generate an improved summary of likes and dislikes using NLTK phrase extraction and frequency analysis"""
        from nltk.tokenize import sent_tokenize, word_tokenize
        from nltk.corpus import stopwords
        from nltk import pos_tag, ne_chunk
        from nltk.chunk import RegexpParser
        from collections import Counter
        import string

        nltk.download('punkt', quiet=True)
        nltk.download('stopwords', quiet=True)
        nltk.download('averaged_perceptron_tagger', quiet=True)

        stop_words = set(stopwords.words('english'))
        sentences = sent_tokenize(text)

        liked_phrases = []
        disliked_phrases = []

        # Define a simple grammar for noun phrase chunking
```



```

grammar = r"""
    NP: {<JJ>*<NN.*>+} # Adjective(s) + Noun(s)
    """

chunk_parser = RegexpParser(grammar)

for sentence in sentences:
    scores = self.sentiment_analyzer.polarity_scores(sentence)
    compound = scores['compound']
    if compound >= 0.05:
        # Extract noun phrases from positive sentences
        words = word_tokenize(sentence)
        words = [w for w in words if w.lower() not in stop_words and w not in string.punctuation]
        tagged = pos_tag(words)
        tree = chunk_parser.parse(tagged)
        for subtree in tree.subtrees(filter=lambda t: t.label() == 'NP'):
            phrase = " ".join(word for word, tag in subtree.leaves())
            liked_phrases.append(phrase)
    elif compound <= -0.05:
        # Extract noun phrases from negative sentences
        words = word_tokenize(sentence)
        words = [w for w in words if w.lower() not in stop_words and w not in string.punctuation]
        tagged = pos_tag(words)
        tree = chunk_parser.parse(tagged)
        for subtree in tree.subtrees(filter=lambda t: t.label() == 'NP'):
            phrase = " ".join(word for word, tag in subtree.leaves())
            disliked_phrases.append(phrase)

    # Count frequency of phrases and get most common ones
    liked_counter = Counter(liked_phrases)
    disliked_counter = Counter(disliked_phrases)

    # Select top 5 phrases or less
    top_liked = [phrase for phrase, count in liked_counter.most_common(5)]
    top_disliked = [phrase for phrase, count in disliked_counter.most_common(5)]

    return {
        'liked': top_liked if top_liked else ['None'],
        'disliked': top_disliked if top_disliked else ['None']
    }

def analyze_feedback(self, text):
    """Complete feedback analysis with sentiment and summary generation"""
    sentiment_result = self.analyze_sentiment(text)
    summary = self.generate_summary(text)

    return {
        'sentiment_analysis': sentiment_result,
        'summary': summary
    }

def main():
    analyzer = CustomerFeedbackAnalyzer()

```

```

feedback = input("Please enter the customer feedback: ")

result = analyzer.analyze_feedback(feedback)

print(f"\nSentiment: {result['sentiment_analysis']['sentiment']} (Score: {result['sentiment_analysis']['score']:.2f})")
print("\nSummary of Customer Feedback:")
print("Liked:", ", ".join(result['summary']['liked']))
print("Disliked:", ", ".join(result['summary']['disliked']))

if __name__ == "__main__":
    main()

```

5.4 Web Interface

The Streamlit interface comprises the following components:

1. **Text Area:** Provides a space for users to input their text for analysis.
2. **“Analyze” Button:** Triggers the sentiment analysis and summarization processes upon user interaction.

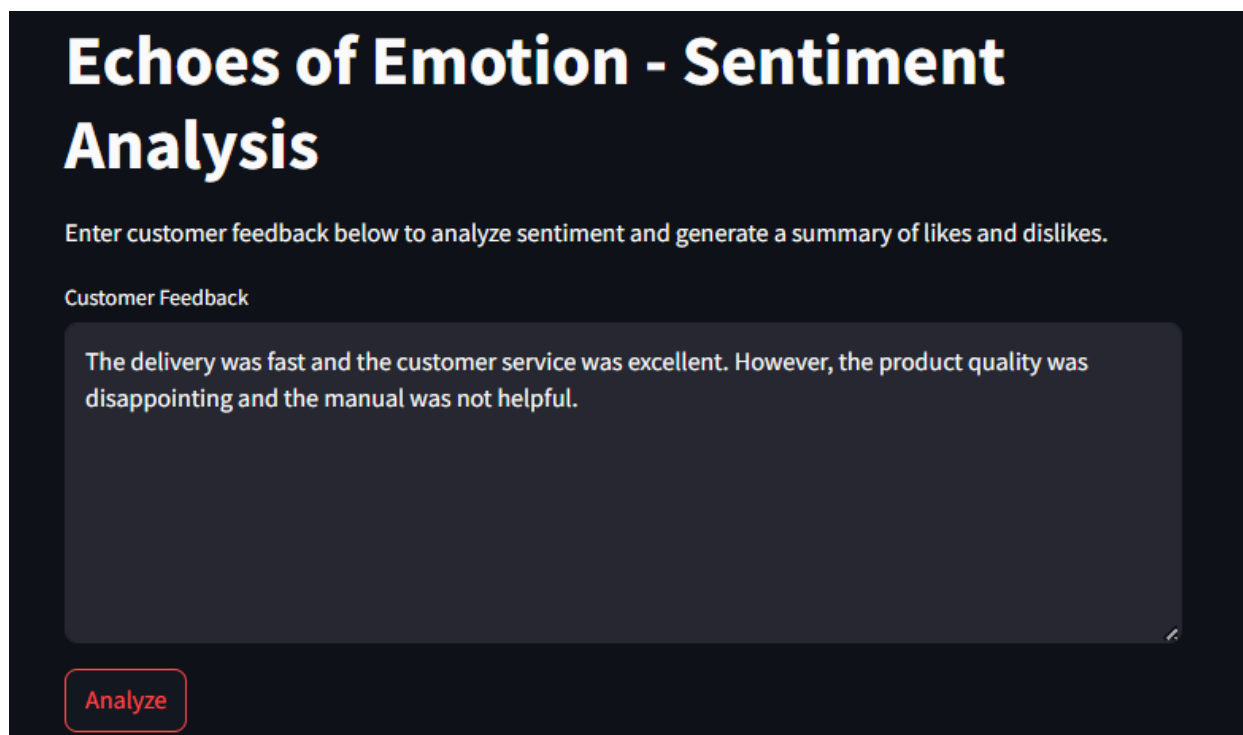


Figure 2: Streamlit Interface

6 Results

The system was tested extensively on the Amazon food reviews dataset, analysing thousands of reviews across various product categories. The key findings are as follows:

- ✚ **Accuracy:** The VADER model demonstrated high precision in classifying sentiment, especially for positive and negative reviews.
- ✚ **Summary Quality:** The summarization module effectively extracted meaningful likes and dislikes, providing valuable insights for business analysis.
- ✚ **Visualizations:** Displays coloured sentiment labels and corresponding emojis to intuitively represent the emotional tone of the input.
- ✚ **Detailed Scores and Summaries:** Presents comprehensive sentiment scores (negative, neutral, positive, compound) alongside generated summaries highlighting key themes.

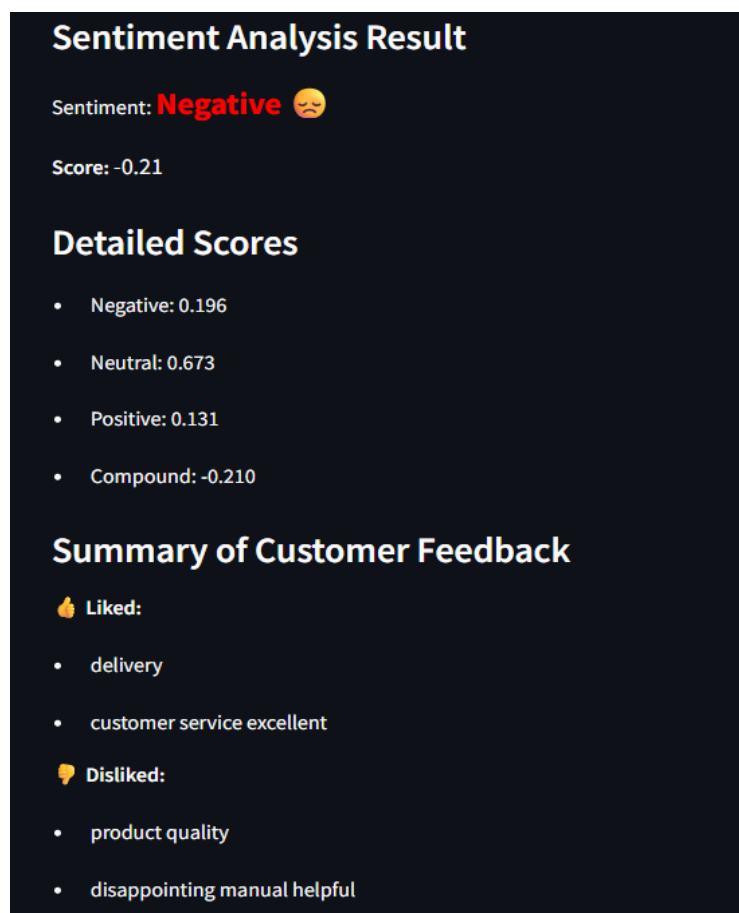


Figure 3: Sample Sentiment Analysis Input

7 Challenges and Solutions

7.1 Challenge: Handling Noisy Data

The reviews contained informal language, slang, and emoticons, complicating sentiment analysis.

Solution: VADER's lexicon, designed specifically for social media text, effectively handled these elements. Additionally, preprocessing steps were applied to remove irrelevant noise.

7.2 Challenge: Summary Accuracy

Extracting precise likes and dislikes from complex and nuanced reviews proved challenging.

Solution: The summarization module was enhanced with part-of-speech (POS) tagging and named entity recognition (NER) to accurately identify key themes.

7.3 Challenge: Interface Responsiveness

Ensuring the Streamlit interface was both user-friendly and responsive across devices was critical.

Solution: Custom CSS styling alongside Streamlit's built-in components optimized the interface for various screen sizes and device types.

7.4 Challenge: Error Handling

Invalid or empty inputs had the potential to disrupt the analysis pipeline.

Solution: Input validation features were added to the Streamlit app, providing users with warnings for empty or invalid inputs.

8 Conclusion

The *Echoes of Emotion* system successfully delivers a robust sentiment analysis platform. By combining VADER's rule-based sentiment model with a Streamlit-powered interactive interface, the system offers an intuitive and effective user experience.

Its comprehensive preprocessing pipeline, accurate sentiment classification, and meaningful summaries make it a valuable tool for both businesses seeking consumer insights and researchers analyzing opinion trends.

Moreover, the system's modular architecture and scalability enable easy adaptation to other datasets and domains. Future enhancements may include integrating machine learning-based sentiment models and expanding support for multilingual text analysis, further broadening the system's applicability.

9 References

1. Hutto, C. J., & Gilbert, E. (2014). VADER: A Parsimonious Rule-based Model for Sentiment Analysis of Social Media Text. *Proceedings of the International AAAI Conference on Web and Social Media*, 8(1).
2. Jurafsky, D., & Martin, J. H. (2023). *Speech and Language Processing* (3rd ed.). Pearson Education.
3. NLTK Documentation (2024). Retrieved from <https://www.nltk.org/>
4. Streamlit Documentation (2024). Retrieved from <https://docs.streamlit.io/>

10 Appendix

10.1 System Requirements

1. **Python:** Version 3.8 or higher
2. **Hardware:** Minimum 4 GB RAM (8 GB recommended), 2 GB free disk space
3. **Internet:** Required to download NLTK mode

10.2 Dependencies

1. nltk (3.8.0 or higher) – for sentiment analysis and NLP preprocessing
2. streamlit (1.20.0 or higher) – for web interface
3. python-dotenv – optional, used for environment variable management if needed

